




**Politecnico
di Torino**

EMERGENCY SIRENS DETECTION





DATASET



PREPROCESSING



MODEL




MQTT



REST



CONCLUSIONS



DATASET


Two labelled classes:

- [Emergency vehicle sirens](#) (932 audio)
- [Traffic noise](#) (902 audio)

Recording at various distance points (10-45 meters)

Different Data Collection method:

- Live Camera
- Online Sources
- Placing Emergency Sirens


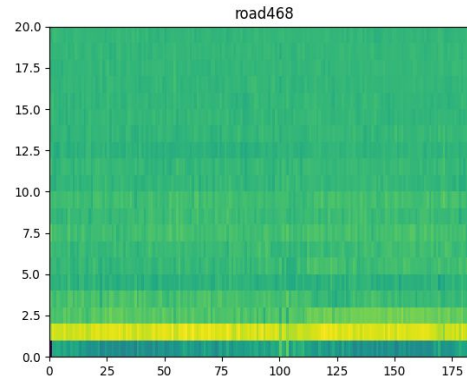
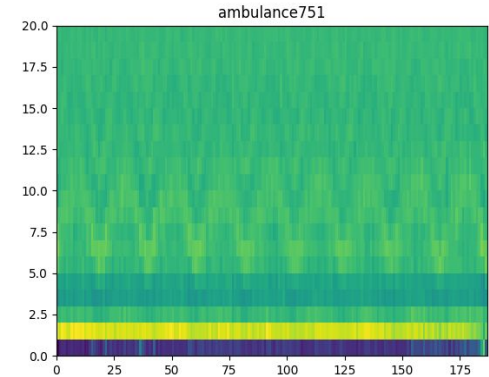
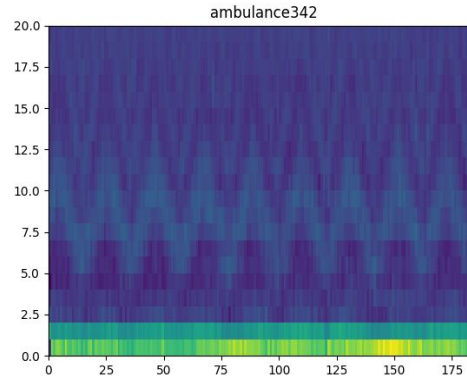



PREPROCESSING

- Audio format converted into WAV
- Up-sampling of the audio to 48 kHz
- Cut audio to 3 seconds
- Padding
- Conversion to MFCC
- Reshape 32x32
- Train/Test/Val split 65/25/10 [%]

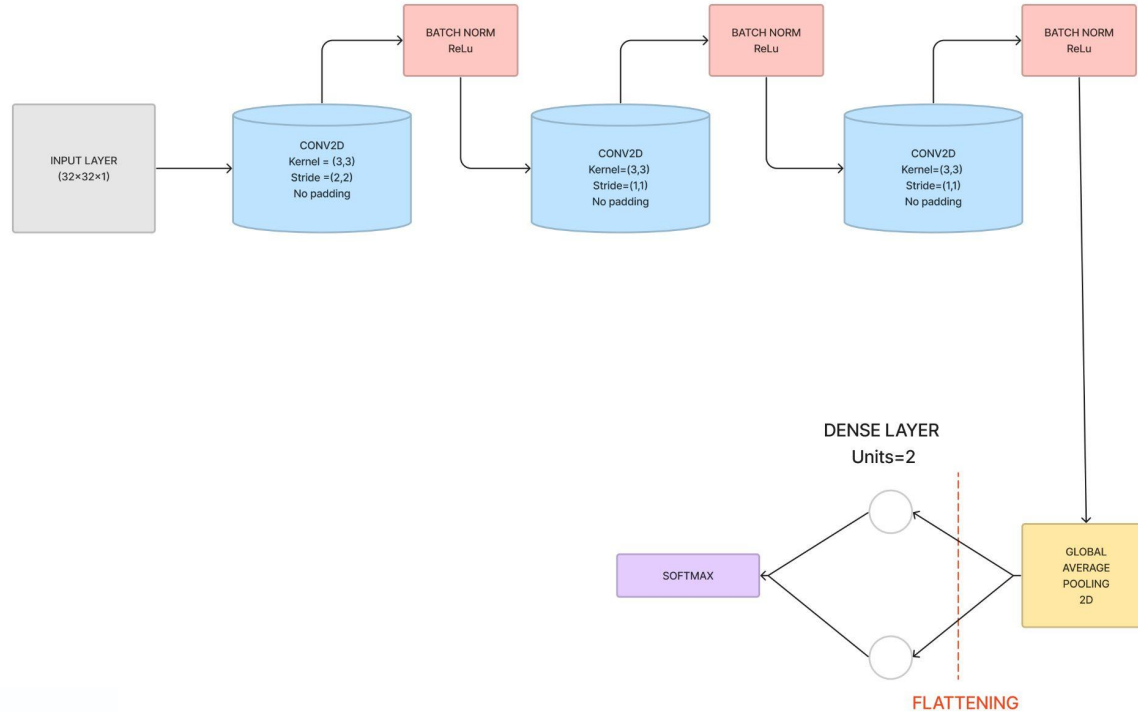
Frame length [s]	0.016
Frame step [s]	0.016
Downsampling rate	48000
Number of mel bins	20

Number of coefficients	20
Lower frequency	20
Upper frequency	24000








MODEL



Batch size	20
Initial learning rate	0.01
End learning rate	1e-5
Epochs	20
Number of filters	64
Pruning alpha	0.2
Initial sparsity	0.2
Final sparsity	0.6



RESULT



Accuracy	95.000%
F1	0.95
Tflite size	31.3 [KB]
Tflite size (.zip)	9.36 [KB]
Median latency	31.4 [ms]
Model parameters	5610 (2798 trainable)

	RoadNoises	EmergencySirens
Precision	0.98	0.92
Recall	0.92	0.98





MQTT

The *publisher* gets the prediction (i.e., it detects an emergency siren) and it sends the MQTT broker the following:

- Sensor name
- Timestamp
- “1”

The *subscriber* stores the received data into a Redis time series. The series are named according to the following rule: “{sensor_name}:emergencySiren”. When a new* {sensor_name} is received, then the related Redis time series is also created.





REST API


The *server* implements three endpoints:


- DeleteSensorHistory
- ListSensors
- SensorHistory

Endpoint	Path	Query
<i>DeleteSensorHistory</i>	optional	no
<i>ListSensors</i>	no	no
<i>Sensor History</i>	optional	optional

The *client* implements the following actions:

- *Geolocation* of a sensor detecting an emergency siren using [folium](#)
- *Computation of statistics* at different granularity levels (sensor or district / year, month, day, day of the week, etc.).
- *Deletion* of time series
- *List* of sensors associated to a time serie





POSSIBLE FUTURE APPLICATION

- Self-driving Car
- Smart Traffic Light
- Road Traffic Monitoring
- Driver Notification





LIMITATIONS AND INTEGRATIONS

Limitations:

- Presence of clacsons
- Multiple Emergency sirens together

Integrations:

- Classification of different emergency sirens
- Increase the dataset using also audio recorded inside the vehicles
- Increment variability in the road noise dataset

THANK YOU FOR YOUR ATTENTION :)

